

CONTACT	Georgia Institute of Technology Atlanta, GA 30332 WeChat: go_bruins Twitter: AhrenJin github.com/Ahren09 (+86)189-1119-7743 (470)962-0241	ahren09.github.io linkedin.com/in/ahren-jin scholar.google.com/citations?user=eY85qm4AAAAJ Semantic Scholar: 2087723977 ORCID: 0000-0002-6974-5970 yjin328@gatech.edu ahren2040@g.ucla.edu
RESEARCH INTERESTS	<ul style="list-style-type: none"> Large Language Models, LLM Agents, Multimodal Learning, Graph Neural Networks, Computational Social Science. 	
EDUCATION	Georgia Institute of Technology Aug. 2022 – May 2027 (Expected) <ul style="list-style-type: none"> Ph.D. Candidate, Computer Science. GPA: 4.0/4.0 University of California, Los Angeles (UCLA) Sep. 2018 – Dec. 2021 <ul style="list-style-type: none"> B.S., Computer Science. GPA: 3.82/4.0. Major GPA: 3.92/4.0. 	
RESEARCH EXPERIENCE	Georgia Tech College of Computing, School of CSE Aug 2022 – Present <i>Graduate Research Assistant</i> Advisor: Dr. Srijan Kumar <ul style="list-style-type: none"> Research Topics: Large Language Models (WebConf'24, ACL'24), LLM Robustness and Safety (Under Review at NAACL'25), Multimodal Models (ACL'24), Recommender Systems and Dynamic Graph Mining (KDD'23), Social Network Analysis (CIKM'24, KDD'23), Fair Graph Mining (CIKM'24). Adobe, REAL Lab May 2024 – Aug 2024 <i>Research Scientist Intern</i> San Jose, CA Advisors: Dr. Gang Wu, Dr. Yu Shen, Dr. Stefano Petrangeli <ul style="list-style-type: none"> Research Topics: Multimodal Large Language Models (MLLMs) Fine-tuning, Web UI and Video Tutorial Understanding. Microsoft Research Asia (MSRA), Social Computing Group Dec. 2020 – July 2022 <i>Research Scientist Intern</i> Beijing, China Advisors: Dr. Xiting Wang, Dr. Jindong Wang, and Dr. Xing Xie <ul style="list-style-type: none"> Research Topics: Large Language Models (EMNLP'24, ICML'24, ICML'23, AAAI'23), LLM Agents (EMNLP'24, ICML'24), Multicultural and Multimodal LLM (Under Review at CVPR'25), Scientometric Analysis (Under Review at NAACL'25), Computational Social Science, Misinformation Detection (KDD'22, AAAI'22), Few-shot Learning (ACL'24, AAAI'23), Explainable AI (AAAI'22). UCLA, Scalable Analytics Institute (ScAi) June 2021 – June 2022 <i>Undergraduate Research Assistant</i> Los Angeles, CA Advisors: Dr. Yizhou Sun and Dr. Wei Wang <ul style="list-style-type: none"> Research Topics: Large Language Models (EMNLP'24), Graph Neural Networks and Data Mining (WWW'23), LLM Fine-tuning (Under Review at KDD'25), Recommender Systems (WWW'23). 	
PUBLICATIONS	In Proceedings <ol style="list-style-type: none"> Yiqiao Jin*, Qinlin Zhao*, Yiyang Wang, Hao Chen, Kaijie Zhu, Yijia Xiao, Jindong Wang. <i>AgentReview: Exploring Peer Review Dynamics with LLM Agents</i>. In Proceedings of EMNLP'24 main track Oral Presentation and SoCal NLP Symposium 2024. Acceptance rate: 20.8%. Yiqiao Jin, Minje Choi, Gaurav Verma, Jindong Wang, Srijan Kumar. <i>MM-Soc: Benchmarking Multimodal Large Language Models in Social Media Platforms</i>. In Proceedings of ACL'24 and the 3rd ALVR Workshop. Acceptance rate: 23.5%. 	

3. **Yiqiao Jin***, Mohit Chandra*, Gaurav Verma, Yibo Hu, Munmun De Choudhury, Srijan Kumar. *Better to Ask in English: Cross-Lingual Evaluation of Large Language Models for Healthcare Queries*. In Proceedings of the ACM Web Conference 2024. **Oral Presentation**.
Acceptance rate: 20.2%.
4. **Yiqiao Jin**, Yeon-Chang Lee, Kartik Sharma, Meng Ye, Karan Sikka, Ajay Divakaran, Srijan Kumar. *Predicting Information Pathways Across Online Communities*. In Proceedings of the 29th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD'23). **Oral Presentation**.
Acceptance rate: 22.1%.
5. **Yiqiao Jin**, Yunsheng Bai, Yanqiao Zhu, Yizhou Sun, Wei Wang. *Code Recommendation for Open Source Project Developers*. In Proceedings of the ACM Web Conference 2023. **Oral Presentation**.
Acceptance rate: 19.2%
6. **Yiqiao Jin**, Xiting Wang, Yaru Hao, Yizhou Sun, Xing Xie. *Prototypical Fine-tuning: Towards Robust Performance Under Varying Data Sizes*. In Proceedings of the 37th AAAI Conference (AAAI'23). **Oral Presentation**.
Acceptance rate: 19.6%.
7. **Yiqiao Jin**, Xiting Wang, Ruichao Yang, Yizhou Sun, Wei Wang, Hao Liao, Xing Xie. *Towards Fine-Grained Reasoning for Fake News Detection*. In Proceedings of the 36th AAAI Conference (AAAI'22). **Oral Presentation**.
Acceptance rate: 14.6%.
8. Qinlin Zhao, Jindong Wang, Yixuan Zhang, **Yiqiao Jin**, Kaijie Zhu, Hao Chen, Xing Xie. *CompeteAI: Understanding the Competition Behaviors in Large Language Model-based Agents*. In Proceedings of the 41st International Conference on Machine Learning (ICML'24), **Oral Presentation**.
Acceptance rate: 27.5%.
9. Yijia Xiao, **Yiqiao Jin**, Yushi Bai, Yue Wu, Xianjun Yang, Xiao Luo, Wenchao Yu, Xujiang Zhao, Yanchi Liu, Haifeng Chen, Wei Wang, Wei Cheng. *Large Language Models Can Be Good Privacy Protection Learners*. In Proceedings of EMNLP'24 main track and SoCal NLP Symposium 2024.
Acceptance rate: 20.8%.
10. Yijia Xiao, Edward Sun, **Yiqiao Jin**, Wei Wang. *RNA-GPT: Multimodal Generative System for RNA Sequence Understanding*. Accepted at NeurIPS'24 MLSB Workshop.
11. Jinghan Zhang, Xiting Wang, **Yiqiao Jin**, Changyu Chen, Xinhao Zhang, Kunpeng Liu. *Prototypical Reward Network for Data-Efficient RLHF*. In Proceedings of ACL'24. **Oral Presentation**.
Acceptance rate: 23.5%.
12. Neng Kai Nigel Neo*, Yeon-Chang Lee*, **Yiqiao Jin**, Sang-Wook Kim, Srijan Kumar. *Towards Fair Graph Anomaly Detection: Problem, New Datasets, and Evaluation*. In Proceedings of the 33rd ACM International Conference on Information and Knowledge Management (CIKM'24).
Acceptance rate: 23.2%.
13. Changyu Chen, Xiting Wang, **Yiqiao Jin**, Victor Ye Dong, Li Dong, Rui Yan, Jim Cao, Yi Liu. *Semi-Offline Reinforcement Learning for Optimized Text Generation*. In Proceedings of the 40th International Conference on Machine Learning (ICML'23).
Acceptance rate: 27.9%.
14. Ruichao Yang, Xiting Wang, **Yiqiao Jin**, Chaozhuo Li, Jianxun Lian, Xing Xie. *Reinforcement Subgraph Reasoning for Fake News Detection*. In Proceedings of the 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD'22).
Acceptance rate: 14.9%.

Preprints

1. **Yiqiao Jin**, Yijia Xiao, Yiyang Wang, Jindong Wang. *Scito2M: A 2 Million, 30-Year Cross-disciplinary Dataset for Temporal Scientometric Analysis*. Under Review at NAACL'25.
2. **Yiqiao Jin**, Andrew Zhao, Yeon-Chang Lee, Meng Ye, Ajay Divakaran, Srijan Kumar. *Empowering Interdisciplinary Insights with Dynamic Graph Embedding Trajectories*.

3. Shudong Liu, **Yiqiao Jin**, Cheng Li, Derek F. Wong, Qingsong Wen, Lichao Sun, Haipeng Chen, Xing Xie, Jindong Wang. *CultureVLM: Characterizing and Improving Cultural Understanding of Vision-Language Models for over 100 Countries*. Under Review at CVPR'25.
4. Sejoon Oh, **Yiqiao Jin**, Megha Sharma, Donghyun Kim, Gaurav Verma, Eric Ma, Srijan Kumar. *LLaVAGuard: Multimodal Safety Guardrails for Jailbreak Attacks on Multimodal Large Language Models*. Under Review at NAACL'25.
5. Yijia Xiao, Edward Sun, **Yiqiao Jin**, Qifan Wang, Wei Wang. *ProteinGPT: Multimodal LLM for Protein Property Prediction and Structure Understanding*. Under Review at KDD'25.
6. Vibhor Agarwal, **Yiqiao Jin**, Mohit Chandra, Munmun De Choudhury, Srijan Kumar, Nishanth Sastry. *MedHalu: Hallucinations in Responses to Healthcare Queries by Large Language Models*. Under Review at NAACL'25.
7. Varun Vankineni, Sai Prasath Suresh, Gaurav Verma, **Yiqiao Jin**, Xiaojing Wang, Rajarajeswari Balasubramaniyan, Srijan Kumar. *Investigating the Effect of Automated vs. Human-Verified Adversarial Samples on Improving the Robustness of LLMs*. Under Review at NAACL'25.
8. Chengyuan Deng, Yiqun Duan, Xin Jin, Heng Chang, Yijun Tian, Han Liu, Henry Peng Zou, **Yiqiao Jin**, Yijia Xiao, Yichen Wang, Shenghao Wu, Zongxing Xie, Kuofeng Gao, Sihong He, Jun Zhuang, Lu Cheng, Haohan Wang. *Deconstructing The Ethics of Large Language Models from Long-standing Issues to New-emerging Dilemmas*.

PROFESSIONAL EXPERIENCE

Amazon, Fulfillment By Amazon (FBA)

June 2020 – Sep. 2020
Seattle, WA

Software Engineer Intern

- Designed and implemented **IAR Manual Analysis**, a scalable and efficient workflow using **AWS Step Functions** and **AWS Lambda**. This service automates the aggregation of data points from multiple sources like **Amazon S3** and **DynamoDB** for **SageMaker** ML model training, handling $\geq 16,000$ requests per summary stage;
- Automated the deployment of the workflow across all AWS Realms (EU/FE/NA) through **CloudFormation**;
- Establish **DataCraft** pipeline to enable automatic data ingestion from **DynamoDB** into the **Andes** dataset catalog, promoting broader internal adoption of these datasets for cross-functional teams and enhancing data accessibility;
- Perform **ablation analysis** on the inventory reconciliation model, identifying key performance bottlenecks and optimizing model performance.

IBM, China Development Laboratories

June 2019 – Sep. 2019
Beijing, China

Software Engineer Intern

- Developed **Compass DataRouter**, a routing service for **Compass** project based on **Golang** and **MongoDB**, reducing memory usage and accelerating data retrieval;
- Enhanced the monitoring dashboard for **Compass** towards a more intuitive and responsive user interface with **React.js**.

SERVICES

Georgia Tech School of Computer Science Graduate Student Association (SCSGSA)

Aug. 2024 –
Atlanta, GA

Communication Chair

- Co-organize the College of Computing (CoC) Graduate Welcome Event with over 1,000 attendees, including new students, faculty, and alumni;
- Lead the design and maintenance of the SCSGSA website (<https://scsgsa.cc.gatech.edu/>);
- Promote SCSGSA events, including workshops, networking sessions, and professional development programs;

ACADEMIC SERVICES

- Area Chair: ICLR'24/23 Tiny Paper
- PC Member/Reviewer
 - **Conferences:**
 - * AAAI'25/24/23;
 - * ACL'24;

- * AISTATS'25/24;
- * ASONAM'24/23;
- * CIKM'24/23;
- * EMNLP'24/23;
- * ICLR'25/24, ICLR'24/23 Tiny Paper;
- * ICML'24;
- * KDD'25/24/23;
- * LoG'24/23/22;
- * NAACL'25;
- * NeurIPS'24/23, NeurIPS'24/23 Datasets & Benchmarks Track;
- * SDM'24;
- * Web Conference'24.

◦ **Journals:**

- * ACM Transactions on Information Systems (TOIS); 2023 –
- * ACM Transactions on Knowledge Discovery from Data (TKDD); 2023 –
- * IEEE Transactions on Knowledge and Data Engineering (TKDE); 2022 –
- * ACM Transactions on Recommender Systems (TORS); 2022 –
- * ACM Transactions on Intelligent Systems and Technology (TIST); 2022 –
- * ACM Transactions on Social Computing (TSC); 2022 –
- * IEEE Transactions on Artificial Intelligence (TAI); 2023 –
- * PLOS ONE; 2023 –
- * International Journal of Data Science and Analytics (JDSA); 2021 –
- * SCIENCE CHINA Information Sciences (SCIS); 2022 –
- * The Journal of Data-centric Machine Learning Research (DMLR). 2023 –

◦ **Workshops**

- * AFT@NeurIPS2023, Ai4Science@NeurIPS2023, Ai4D3@NeurIPS2023, GenBio@NeurIPS2023, TGL Workshop@NeurIPS2023;
- * SPIGM@ICML2023.

● Volunteer: VLDB'24

HONORS AND
AWARDS

- AAAI Student Scholarship. 2022, 2023
- Microsoft Research “Star of Tomorrow” Award of Excellence. 2021
- UCLA Dean’s Honor List for Superior Academic Achievement. 2019 – 2021
 - 5 times: Spring 2019, Winter 2020, Spring 2020, Winter 2021, Spring 2021